

5-Day Sprint to 80% Functional

Real deadline: **Tuesday, 1 September 2026** — 5 working days from today (Thu 27 Aug). This replaces the earlier Sept-20 / December-paced plan; that timeline no longer applies here.

5 working days left

Target: 80% functional

Deadline: Tue 1 Sep 2026

6 members, 2 powerful laptops

01 What “80% functional” means — say the bar out loud

Building the full architecture as originally spec'd — in-browser vision model, full 3-tier privacy guard, PVM memory, cross-browser support, full benchmark suite — in 5 days with 6 people is not realistic, and pretending otherwise is how a team ends up with six unfinished pieces on Sept 1 instead of one working one. The honest move is to cut scope on purpose, not by accident.

In scope — this is the 80%, must work live on the mock site:

- Extension loads and captures DOM state of the mock test site.
- Tier-1 DOM-rule PII detection + deterministic redaction (blackout/mask) runs before any network call.
- Sanitized payload goes to a **local** LLM, self-hosted via Ollama on Server AI's own GPU; server returns one structured action (click/type/scroll) via a frozen JSON schema.
- Extension validates and executes that action locally.
- Level-1 deterministic verification (URL/DOM check) confirms the action actually worked.
- One canary-value test proves zero raw PII reached the server — the privacy proof.
- A basic privacy inspector panel shows what was redacted and what was sent.
- One complete task runs start-to-finish, repeatably, on the mock site.

Deferred past Sept 1 — the hardest 20%, cut on purpose:

- The in-browser local vision model (ONNX Runtime Web + WebGPU) — normally the multi-week piece on its own. Perception is **DOM/A11y-only** for this sprint.
- PVM verified-memory caching.
- Multi-site testing and Firefox compatibility.
- The full 5-metric benchmark suite.
- Visual/face redaction — text and DOM-field PII only for now.

02 Hardware — reasoning moved onto Server AI's GPU

Deferring the local vision model means Perception doesn't need standing GPU priority this sprint. But this project runs its reasoning model locally too, not via a cloud API — a deliberate choice, not a privacy requirement (the redaction step is what protects PII; see *docs/ARCHITECTURE.md*). That means **Server AI's laptop is the standing GPU host all week**, running Ollama continuously. **Integration** gets the second machine for wiring and the golden-path stress test. Every other role below is scoped to not need a laptop with a GPU at all — the server defaults to a stub with zero setup.

03 Day-by-day, all six roles

Role	Fri 28 Aug	Sat 29 Aug	Sun 30 Aug	Mon 31 Aug	Tue 1 Sep
Extension	MV3 skeleton + content script into mock site	Action executor (click/type/scroll) + validator	type_secret pattern + error handling	Full-team integration: fix breakages live	Bug triage only — no new features
Perception	DOM/A11y capture, normalize page state	Hand off state to Privacy Guard + Server AI	Polish capture on mock site's specific fields	Support integration + stress-test capture	Bug triage only

Role	Fri 28 Aug	Sat 29 Aug	Sun 30 Aug	Mon 31 Aug	Tue 1 Sep
Privacy Guard	Tier-1 DOM-rule detection + deterministic redaction	Redaction validator + sanitized payload builder	Canary-test harness + Privacy Firewall block-on-fail	Run canary test, capture SHA-256 proof	Bug triage only
Server AI	FastAPI skeleton; freeze action schema w/ Extension; install Ollama + pull model	Sanitized-context endpoint live; hardcoded action to unblock	Real local LLM (Ollama) reasoning returns valid structured actions	Support integration + canary proof from server side	Bug triage only
PVM/Verify	Stub only — help Integration build the mock site	Level-1 deterministic verification (URL/DOM check)	Basic recovery loop for failed actions	Stress-test the one golden task for reliability	Bug triage only
Integration	Build the mock test site — ships today, blocks everyone	First end-to-end smoke test, however broken	Privacy inspector / dashboard v1	Lock demo script + record a backup video	Final PPT + submit before deadline

Day 4 is a mandatory full-team sync — this is where six parallel 3-day builds either click together or reveal a broken assumption while there's still one day to fix it. Don't skip it to keep coding.

1

Extension & Execution

Owens this sprint: Manifest V3 skeleton, content scripts, background service worker, action executor, client action validator.

Role 1 of 6 · EXTENSION · 5-day sprint

HW Hardware

None needed. Any laptop, all 5 days.

DAY Day-by-day

Fri 28 Aug

- Set up the extension/ skeleton — manifest.json, content script, background service-worker stub.
- Get a “hello world” content script injecting into the mock site the moment Integration has it up.
- Agree the message-passing contract with Perception and Privacy Guard.

Sat 29 Aug

- Build the action executor: map click/type/scroll to real browser APIs.
- Build the client action validator — target exists, action type allowed, confidence threshold met.

Sun 30 Aug

- Build the type_secret pattern — server requests a semantic op, real value resolved locally.
- Add basic error handling so a bad server response doesn't crash the extension mid-demo.

Mon 31 Aug

- Full-team integration sync — plug into the live pipeline and fix whatever breaks.
- Help PVM/Verification wire the recovery path if actions fail.

Tue 1 Sep

- Bug triage only. No new features today — stability over scope.

DEP Dependencies

YOU DEPEND ON	Server AI (action schema), Privacy Guard (nothing blocks you Day 1–2).
DEPENDS ON YOU	Everyone — you're the runtime host. If Day 1's skeleton slips, nobody else has anything to plug into.

2

Perception & Local Inference

Owens this sprint: DOM/A11y capture and page-state normalization for this sprint. Local vision model (ONNX/WebGPU) is deferred past Sept 1.

Role 2 of 6 · PERCEPTION · 5-day sprint

HW Hardware

None needed this sprint — DOM-only scope means no GPU work until after the deadline.

DAY Day-by-day

Fri 28 Aug

- DOM/accessibility capture: element roles, labels, ARIA, form types.
- Normalize captured state into one structured representation everyone else can consume.

Sat 29 Aug

- Hand off the normalized state format to Privacy Guard (for detection) and Server AI (for the sanitized payload).

Sun 30 Aug

- Polish capture against the mock site's actual fields — this is where real bugs hide, not in generic DOM parsing.

Mon 31 Aug

- Full-team integration sync.
- Stress-test capture on the golden-path task, repeatably.

Tue 1 Sep

- Bug triage only.

DEP Dependencies

YOU DEPEND ON	Extension (content-script scaffold), Integration (mock site).
DEPENDS ON YOU	Privacy Guard and Server AI both consume your normalized state directly — a format change after Day 1 breaks both of them.

3

Privacy Guard & Redaction

***Owens this sprint:** Tier-1 DOM-rule PII detection, deterministic redaction, redaction validation, the canary-value proof.*

Role 3 of 6 · PRIVACY · 5-day sprint

HW Hardware

None needed. Any laptop, all 5 days.

DAY Day-by-day

Fri 28 Aug

- Tier-1 DOM-rule detection only for this sprint (input type/autocomplete/labels/ARIA) — skip the text-NER and visual tiers for now.
- Deterministic redaction renderer: blackout/mask before anything is sent.

Sat 29 Aug

- Redaction validator — confirm every flagged field is actually handled.
- Build the sanitized-payload builder that Server AI's endpoint will consume.

Sun 30 Aug

- Design and build the canary-value test harness (e.g. CANARY_EMAIL_12345).
- Privacy Firewall: block the request outright on validation failure — this is your highest-value single feature for the demo.

Mon 31 Aug

- Full-team integration sync.
- Run the canary test end-to-end with Server AI; capture the proof artifact (payload log / hash).

Tue 1 Sep

- Bug triage only.

DEP Dependencies

YOU DEPEND ON	Perception (DOM state to run detection on), Extension (network-intercept point).
DEPENDS ON YOU	Server AI must only ever receive your sanitized output — this is the single most judge-visible claim in the whole project.

4

Server AI

Owens this sprint: FastAPI + Pydantic, the local Ollama reasoning call, the frozen action schema, context validation.

Role 4 of 6 · SERVER · 5-day sprint

HW Hardware

Heavy — this is the standing GPU host all week. Your own laptop runs Ollama continuously; keep it charged, powered on, and reachable on the network for anyone testing the real reasoning path.

DAY Day-by-day

Fri 28 Aug

- FastAPI skeleton with a sanitized-context endpoint.
- Freeze the structured action schema (click/type/scroll + type_secret) with Extension **today** — this blocks their Day 2 work.
- Install Ollama and pull the chosen model (~5GB download) so it's ready before you need it.

Sat 29 Aug

- Get the endpoint live end-to-end using the stub backend first, even if it returns a hardcoded action — unblocking Extension matters more than real reasoning on Day 2.

Sun 30 Aug

- Switch to the local Ollama backend (REASONING_BACKEND=ollama). Get real reasoning over the sanitized context returning valid, schema-conformant actions — and a safe fallback when the model returns malformed JSON.

Mon 31 Aug

- Full-team integration sync.
- Run the canary proof from the server side — confirm the request log shows zero raw PII received, regardless of which backend is active.

Tue 1 Sep

- Bug triage only.

DEP Dependencies

YOU DEPEND ON	Privacy Guard (sanitized context format), Extension (the action schema, agreed Day 1).
DEPENDS ON YOU	Extension needs your response format to build its validator. Everyone testing real (non-stub) reasoning depends on your laptop being on and reachable — say so in the team chat if you need to step away.

5 PVM & Verification

Owens this sprint: For this sprint: Level-1 deterministic verification and a basic recovery loop. Full PVM memory/caching is deferred past Sept 1.

Role 5 of 6 · PVM · 5-day sprint

HW Hardware

None needed. Any laptop, all 5 days.

DAY Day-by-day

Fri 28 Aug

- No verification target exists yet — spend today helping Integration build the mock site instead.

Sat 29 Aug

- Build Level-1 deterministic verification only: did the URL/DOM change the way the action predicted?

Sun 30 Aug

- Basic recovery loop: on a failed/ambiguous action, capture new state and retry once, safely.

Mon 31 Aug

- Full-team integration sync.
- Stress-test the one golden-path task repeatedly — this is where reliability either holds or doesn't.

Tue 1 Sep

- Bug triage only.

DEP Dependencies

YOU DEPEND ON	Extension (action + result state to verify), Server AI (a real action to verify against).
DEPENDS ON YOU	Extension's recovery path hooks into your verification result.

6

Integration, Benchmark & Demo

Owens this sprint: *The mock test site, end-to-end wiring, the privacy inspector v1, the demo script — full benchmark suite is deferred past Sept 1.*

Role 6 of 6 · INTEGRATION · 5-day sprint

HW Hardware

Second-claim priority on a laptop for wiring/testing sessions.

DAY Day-by-day

Fri 28 Aug

- Build the mock e-commerce/booking test site and ship it **today** — every other role is blocked without it.

Sat 29 Aug

- Run the first end-to-end smoke test across whatever pieces exist, however broken — the goal is finding the real integration gaps early.

Sun 30 Aug

- Build the privacy inspector / dashboard v1: redacted items, allow/block status. No sensitive values ever shown.

Mon 31 Aug

- Run and lead the full-team integration sync.
- Lock the demo script; record a backup video in case the live demo fails.

Tue 1 Sep

- Final PPT with real screenshots from the working prototype, not mockups. Submit before the deadline.

DEP Dependencies

YOU DEPEND ON	Every other role landing their Day 1–3 pieces on schedule.
DEPENDS ON YOU	Everyone — the mock site is the Day-1 blocker, and Day 4's sync is where six separate builds become one working demo.